

	HH_type	Religion	Social_Group	Regular_salary_earner	Possess_ration_card	Sex	Age	Marital_Status	Education	Meals_At_Home	I
0	2.0	1.0	3.0	1.0	1.0	1	50	2.0	8.0	59.0	
1	2.0	3.0	9.0	1.0	1.0	2	40	3.0	12.0	56.0	
2	1.0	1.0	9.0	1.0	1.0	1	45	2.0	7.0	60.0	
3	2.0	3.0	9.0	1.0	1.0	1	75	3.0	6.0	60.0	
4	1.0	1.0	9.0	2.0	1.0	1	30	2.0	7.0	59.0	

```
# Impute missing values with mean
imputer = SimpleImputer(strategy='mean')
X_imputed = imputer.fit_transform(X)
X_imputed
```

```
array([[2.0000e+00, 1.0000e+00, 3.0000e+00, ..., 5.0000e+00, 4.7510e+04,
        4.1100e+02],
       [2.0000e+00, 3.0000e+00, 9.0000e+00, ..., 2.0000e+00, 8.5102e+04,
        3.3100e+02],
       [1.0000e+00, 1.0000e+00, 9.0000e+00, ..., 5.0000e+00, 4.9219e+04,
        1.2100e+02],
       ...,
       [5.0000e+00, 1.0000e+00, 9.0000e+00, ..., 7.0000e+00, 4.1001e+04,
        9.2000e+02],
       [2.0000e+00, 1.0000e+00, 9.0000e+00, ..., 5.0000e+00, 4.7211e+04,
        5.2200e+02],
       [1.0000e+00, 1.0000e+00, 9.0000e+00, ..., 7.0000e+00, 1.1130e+03,
        6.1100e+02]])
```

```
# Ensure 'y' is a binary factor
y = y.astype('category')
```

```
# Scale numeric variables if needed (optional)
X_imputed_scaled = StandardScaler().fit_transform(X_imputed)
X_imputed_scaled
```

```
array([[ -0.33833318, -0.4138791 , -0.47257591, ...,  0.18974726,
         0.28981662, -0.70113625],
       [-0.33833318,  1.34847893,  1.42994352, ..., -1.14520676,
         1.58409885, -1.02440146],
       [-0.80105208, -0.4138791 ,  1.42994352, ...,  0.18974726,
         0.34865702, -1.87297265],
       ...,
       [ 1.04982353, -0.4138791 ,  1.42994352, ...,  1.0797166 ,
         0.06571354,  1.35563868],
       [-0.33833318, -0.4138791 ,  1.42994352, ...,  0.18974726,
         0.27952214, -0.25260576],
       [-0.80105208, -0.4138791 ,  1.42994352, ...,  1.0797166 ,
        -1.30761934,  0.10702679]])
```

```
# Fit the probit regression model
probit_model = Probit(y, X_imputed_scaled)
probit_result = probit_model.fit(maxiter=1000)
```

```
# Print model summary
print(probit_result.summary())
```

```
Optimization terminated successfully.
Current function value: 0.663551
Iterations 5
```

```

                Probit Regression Results
=====
Dep. Variable:          non_veg      No. Observations:          101662
Model:                  Probit      Df Residuals:              101648
Method:                  MLE        Df Model:                  13
Date:                   Mon, 01 Jul 2024  Pseudo R-squ.:          -0.05190
Time:                   14:18:08      Log-Likelihood:          -67458.
converged:              True         LL-Null:                -64129.
Covariance Type:        nonrobust     LLR p-value:             1.000
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
x1             -0.0285      0.004      -6.346      0.000      -0.037      -0.020
x2              0.1579      0.005      34.313      0.000      0.149      0.167
x3             -0.1254      0.004     -30.423      0.000      -0.133      -0.117
x4             -0.0135      0.005      -2.880      0.004      -0.023      -0.004
x5             -0.0198      0.004      -4.674      0.000      -0.028      -0.011
x6             -0.0475      0.005      -9.236      0.000      -0.058      -0.037
x7             -0.0176      0.005      -3.839      0.000      -0.027      -0.009
x8              0.0478      0.005       9.059      0.000      0.037      0.058
x9             -0.0322      0.005      -6.551      0.000      -0.042      -0.023
x10            0.1669      0.004      39.700      0.000      0.159      0.175
=====
```

x11	-0.0968	0.004	-23.621	0.000	-0.105	-0.089
x12	0.0200	0.004	4.613	0.000	0.011	0.028
x13	0.0756	0.005	15.497	0.000	0.066	0.085
x14	0.0398	0.005	8.239	0.000	0.030	0.049

=====

```
# Predict probabilities
predicted_probs = probit_result.predict(X_imputed_scaled)

# Convert probabilities to binary predictions using a threshold of 0.5
predicted_classes = np.where(predicted_probs > 0.5, 1, 0)

# Confusion Matrix
cm = confusion_matrix(y, predicted_classes)
print("Confusion Matrix:")
print(cm)

# Classification Report
print("Classification Report:")
print(classification_report(y, predicted_classes))
```

```
→ Confusion Matrix:
[[22233 10839]
 [29282 39308]]
Classification Report:
              precision    recall  f1-score   support


     0       0.43        0.67        0.53       33072
     1       0.78        0.57        0.66       68590

 accuracy          0.61          0.61          0.61       101662
 macro avg         0.61          0.62          0.59       101662
 weighted avg      0.67          0.61          0.62       101662
```

```
# Calculate ROC curve and AUC
fpr, tpr, thresholds = roc_curve(y, predicted_probs)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure()
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc="lower right")
plt.show()

print(f"AUC: {roc_auc:.2f}")
```



AUC: 0.66

Start coding or [generate](#) with AI.